


The Big Picture

What we build, how our architecture is structured, and why production-grade AI is different from a demo. Use this module to frame any technical conversation with a prospect.

Our Enterprise AI Architecture

Every engagement operates within a five-layer model. Understanding these layers lets you explain exactly where our work sits inside a client's existing technology estate — and why every layer matters.



The Four AI Operating Modes

Predictive DATA-DRIVEN

Uses historical and real-time data to forecast outcomes before they happen. Client examples include inventory demand forecasting, fraud detection scoring, credit risk assessment, and workforce capacity planning.

ML models time-series scoring APIs

Conversational USER-FACING

AI-powered interfaces that respond to natural language. Examples: internal knowledge assistants, customer support bots, HR policy Q&A, and procurement request handling embedded in Slack or a web portal.

LLMs RAG chat APIs

Generative CONTENT

AI that creates new documents or structured content at scale. Examples: automated report generation, loan application summaries, contract drafting, compliance document review, and email response drafting.

LLMs templates structured output

Analytical INSIGHT

AI that surfaces hidden patterns in operational data. Examples: customer churn signals, pipeline health scoring, document classification, and sentiment analysis on support tickets to route them faster.

NLP classification embeddings

Production-Grade vs. Proof-of-Concept

Sales tip: Most prospects have already seen an AI demo. The real question they're asking is: "Can this actually work inside our environment, with our data, under our security policies?" This table is your answer.

DIMENSION	PROOF-OF-CONCEPT (WHAT THEY'VE LIKELY SEEN)	WE PRODUCTION-GRADE
Data	Sample CSV or synthetic data	✓ Live integrations with real CRM / ERP
Security	No auth, shared API keys	✓ Role-based access, audit logs, isolated tenants
Reliability	Demo environment only, no SLA	✓ Uptime monitoring, error handling, fallback logic
Monitoring	None	✓ Drift detection, evaluation pipelines, alerting
Integration	Standalone app, copy-paste workflow	✓ Embedded in existing tools and workflows
Compliance	Not considered in the PoC	✓ Designed for regulated industries from day one
Ownership	Vendor lock-in, hosted platform	✓ Deployed in your environment — you own the system

MODULE 02

Services Deep Dive

A technical breakdown of each of our services — how it works under the hood, what systems it connects to, and what the client receives.


🕒 AI Workflow Audit

What it is

A structured analysis of one specific business process to identify the highest-value AI automation entry point. Delivered as a scoped report with a concrete, costed implementation pathway.

What the client receives

- ✓ **Process map** annotated with bottleneck markers
- ✓ **Automation score** per process step (effort vs. impact)
- ✓ **Top use case** with rough ROI estimate
- ✓ **Implementation roadmap** phased into 3 stages



🕒 Internal Knowledge Assistant


Key technical concept — RAG: This service uses Retrieval-Augmented Generation. Instead of the AI "guessing" answers from its training data, it first searches the client's actual documents, retrieves the most relevant sections, then generates an answer grounded only in those specific documents — with citations.

How RAG works (for the sales conversation)

Think of it like giving the AI a private search engine over your company's documents. When someone asks a question, the AI searches first, finds the right policy or procedure, then answers — quoting the source. It never makes up an answer that isn't in your documents.

What gets connected

- ✓ **Confluence / Notion** — internal wikis and runbooks
- ✓ **SharePoint / Google Drive** — file storage and policies
- ✓ **PDF / Word documents** — procedures and compliance docs
- ✓ **Slack history** — surfacing institutional knowledge




🕒 Document Automation

What it automates

Repetitive document processing: extracting structured data from forms, classifying incoming document types, generating outputs from unstructured text, and routing documents to the correct downstream workflow — without manual handling.

Human-in-the-loop (a feature, not a limitation)

When the AI's confidence score falls below the configured threshold, the document is flagged for human review rather than processed automatically. This means the system never silently makes a high-stakes mistake.




🕒 Customer Support Assistant

How intelligent routing works

The AI classifies each incoming message by intent (billing query, technical issue, complaint, general inquiry) and assigns a confidence score. High-confidence intents are handled automatically; ambiguous ones escalate to a human agent — with full conversation context pre-populated.

CRM integration behavior

- ✓ **Reads** full customer history before composing a reply
- ✓ **Logs** every AI interaction to the CRM record
- ✓ **Creates** tickets or updates case records automatically
- ✓ **Escalates** with conversation context pre-loaded for the agent




🕒 Reporting & Research Copilot

What it replaces

Hours of manual data gathering, spreadsheet consolidation, and report drafting. The AI connects directly to data sources, queries them on schedule, summarizes findings in your preferred format, and delivers the report via email or Slack — without human intervention.

Data connectors supported

- ✓ **SQL databases** — direct query access via read-only credentials
- ✓ **BI tools** — Tableau, Power BI, Looker data exports
- ✓ **APIs** — internal and third-party data feeds
- ✓ **Spreadsheets** — Google Sheets, Excel via scheduled sync



🕒 AI Implementation Advisory

What makes this different from generic strategy advice


Recommendations are scoped to the client's actual tech stack, team capabilities, and data readiness. The output is a prioritized roadmap with effort/impact scoring — not a slide deck full of buzzwords. Every recommendation is tied to a specific, implementable next action.

Assessment dimensions

- ✓ **Data readiness** — quality, availability, and governance status
- ✓ **Integration feasibility** — existing systems and API availability
- ✓ **Team capability** — skill gaps and training requirements
- ✓ **ROI modeling** — time and cost savings estimated per use case

How client data flows from existing source systems through our AI pipelines — and back into the business. Use this to answer "where does our data go?"

Data Integration Map



Deployment Models

CLOUD-HOSTED

Deployed within the client's own AWS, GCP, or Azure account. Data never leaves their cloud environment. We manage the deployment — not the data. Most common for mid-market clients.

ON-PREMISES

Fully air-gapped within the client's physical or virtual infrastructure. Uses open-source models (Llama 3, Mistral) running locally. Required for highly regulated environments.

HYBRID

Sensitive regulated data stays on-premises; non-sensitive processing uses cloud inference. Common in lending — some datasets are regulated, others aren't.

Data Isolation & Ownership

Key assurance for enterprise procurement: Each client deployment is fully isolated. No shared model learns from one client's data and applies it to another. Your data is never used to train general-purpose models.

- ✓ Separate database schema or dedicated instance per client
- ✓ API keys and credentials scoped exclusively to each deployment
- ✓ Audit logs are client-specific and fully exportable on request
- ✓ Data retention policies configurable per engagement agreement
- ✓ Client retains full ownership — system lives in their environment

MODULE 04

Security & Compliance

How we handle data security, access control, audit trails, and regulated industry requirements. Use this for enterprise procurement and legal review conversations.

Security Boundary Map




Fig 4.1 — Security boundaries. Sensitive credentials exist only server-side. Client never touches secrets directly.

Security Measures in Plain English

Access Control

- ✓ **Role-based permissions** — users see only what they need
- ✓ **Service accounts** — systems talk with scoped keys only
- ✓ **Row-level security** enforced at database layer (RLS)
- ✓ **Admin access** server-side only, never client-side

Audit & Traceability

- ✓ **Every AI action logged** — who, when, what input, what output
- ✓ **Immutable records** — logs cannot be retroactively altered
- ✓ **Export available** — full audit trail exportable for compliance review
- ✓ **Timestamps** — created_at / updated_at on all records

Input Security

- ✓ **Server-side validation** — all input sanitized before processing
- ✓ **Length limits** — prevent injection and overflow attempts
- ✓ **Honeypot fields** — bot and spam detection on all endpoints
- ✓ **Domain validation** — rejects personal emails for B2B flows

AI-Specific Security

- ✓ **Prompt injection mitigation** — user input sanitized before LLM
- ✓ **Output filtering** — AI responses screened before delivery
- ✓ **Domain guardrails** — model constrained to defined scope
- ✓ **Confidence thresholds** — low-confidence outputs escalate to humans

Regulatory Compliance Context

Sales guidance: We do not currently hold SOC 2 or ISO 27001 certifications. When a prospect asks, be honest and pivot to architecture: "We're designed with these frameworks in mind and can provide full documentation for your security review team."

REGULATION	WHO IT APPLIES TO	HOW WE ADDRESS IT
GDPR	EU-based clients or clients handling EU personal data	Data stays in client-controlled cloud; data processing agreements available on request
PDPO	Hong Kong personal data processing	Primary jurisdiction for we; data residency options built into HK deployments
MAS TRM	Singapore financial institutions	Audit trail requirements met; model risk documentation available
HKMA	HK banks and licensed lenders	On-premises deployment available for regulated data; full audit exports
SOC 2	US enterprise and SaaS buyers	Architecture follows SOC 2 principles; formal certification on roadmap; security questionnaire available

MODULE 05

MLOps & Monitoring

How we ensure AI systems keep working correctly after they go live. Use this to answer "what happens after deployment?" — the question every technical buyer asks eventually.




Fig 5.1 — MLOps lifecycle: Deploy → Monitor → Evaluate → Retrain → back to Deploy.

What Each Stage Means for the Client

Monitor — What is tracked

- ✓ **Response latency** — is the AI responding within SLA?
- ✓ **Error rates** — are requests failing or timing out?
- ✓ **Input distribution** — are query types shifting?
- ✓ **Output consistency** — are answers staying on-topic?

Evaluate — What is scored

- ✓ **Accuracy** — are answers factually correct?
- ✓ **Relevance** — are answers appropriate for each query?
- ✓ **Drift score** — how far has performance shifted from baseline?
- ✓ **Human feedback** — thumbs up/down signals from end users

Model Drift — Explained for Clients

What is drift? An AI model trained on last year's data may become less accurate as the world changes. A lending risk model trained before an interest rate hike may score applications differently than intended after the hike. Drift monitoring catches this before it causes problems.


Data drift

Incoming data starts looking different from the training data. Example: new product categories appear in queries that the model wasn't trained on. Detected by comparing statistical distributions of inputs over time.

Concept drift

The relationship between input and correct output changes. Example: customer sentiment language evolves and the model no longer interprets it correctly. Detected by declining accuracy scores on sampled outputs.

Human-in-the-Loop Escalation



Why this matters in sales: Human-in-the-loop is a professional risk management feature. It means the AI never silently delivers a wrong answer on a high-stakes decision. Frame it as governance, not a limitation.


The infrastructure we run on — demonstrating that we operate with the same engineering standards we build for clients. Use this to establish technical credibility.

Technology Stack

LAYER	TECHNOLOGY	VERSION	PURPOSE
Frontend Framework	Next.js + React	16.2 / 19.2	Server-side rendering, App Router, API routes — modern web architecture
Language	TypeScript	5.x strict	Type-safe development — catches errors at compile time, not runtime
Styling	Tailwind CSS	4.x	Utility-first design system — consistent, responsive, maintainable UI
Animation	Framer Motion	11.x	Scroll-triggered reveals, timeline animations, reduced-motion support
Database	Supabase (PostgreSQL)	Latest	Client portal data, lead capture, project milestones — with row-level security
Email delivery	Resend + Web3Forms	—	Multi-tier fallback: Resend (primary) → Web3Forms (browser fallback) → mailto
Hosting	Netlify	—	Serverless deployment with Next.js plugin — global CDN, automatic deploys
Runtime	Node.js	20 LTS	Server-side API route execution — current long-term support version
Internationalisation	Custom i18n context	—	English + Traditional Chinese (zh-HK) — purpose-built for HK & international clients

Client Portal Architecture

Each client receives a unique reference code that unlocks their project portal — showing current stage, milestones, and progress in real time, backed by Supabase with a Task Manager fallback for resilience.



Why This Matters in Sales

Credibility point: When a prospect asks "have you built production systems like this before?" — walk them through our own infrastructure. We don't just consult on production AI systems; we run one ourselves. A bilingual (English / Traditional Chinese) client portal, backed by a production PostgreSQL database with proper access controls and service redundancy, is a working demonstration of our engineering standards.

MODULE 07

Sales FAQ

The 15 most common technical questions from enterprise prospects — with answers your sales team can use immediately. Click any question to expand.

Where is our data stored?

Your data stays in your own environment. For cloud deployments, we deploy into your existing AWS, GCP, or Azure account — we never have persistent access to your data. For on-premises deployments, nothing leaves your physical network. We do not store client data in our own infrastructure.

Can this connect to our existing CRM or ERP?

Yes. We build native integrations with Salesforce, HubSpot, SAP, Oracle, and most SQL databases via standard APIs and connectors. Integration scope is agreed upfront during the AI Implementation Advisory phase so there are no surprises mid-project. We use read-only access initially and expand only as needed for each use case.

What LLM do you use — is it OpenAI / ChatGPT?

We are model-agnostic and select the best option for your use case and compliance requirements. Options include OpenAI GPT-4o, Anthropic Claude, Google Gemini, or open-source models like Llama 3 and Mistral for fully on-premises deployments. For regulated clients, we default to open-source models that run entirely within your environment — zero data leaves your network.

Who has access to our data?

Access is scoped by role and enforced at both application and database level. Only the service accounts required for each specific integration have access. No engineers on our team hold standing access to production client data — any deployment access is time-limited and logged in full. Audit logs are available for your security team to review at any time.

What happens if the AI gives a wrong answer?

Every deployment includes a human-in-the-loop escalation path. When the AI's confidence score falls below the threshold agreed for your use case, it flags the item for human review rather than delivering a low-confidence answer automatically. All AI outputs are also logged, so any errors can be traced, corrected, and used to improve the model through the retraining cycle.

Is this GDPR / PDPO compliant?

Our architecture is designed to support compliance with GDPR and Hong Kong's PDPO. Data stays in client-controlled environments, data processing agreements are available on request, data retention is configurable per agreement, and audit trails are fully exportable. We recommend your Data Protection Officer review our architecture documentation as part of procurement — we can provide this on request.

Can we run this on-premises?

Yes. For clients in regulated industries (banking, insurance, healthcare) or with strict data residency requirements, we deploy fully air-gapped solutions using open-source models such as Llama 3 or Mistral. The AI never calls an external API — all inference runs locally on your infrastructure. This requires appropriate hardware but eliminates any external data exposure entirely.

What is the typical integration timeline?

Timelines depend on the service and data source complexity. An AI Workflow Audit delivers in 1–2 weeks. A fully deployed Internal Knowledge Assistant typically takes 4–8 weeks end-to-end, including data source access approvals. The AI Implementation Advisory phase (2 weeks) runs first and produces a detailed timeline specific to your stack, constraints, and access approval process.

How do we know it's working correctly after go-live?

Every production deployment includes monitoring dashboards showing AI performance metrics, accuracy scores, and drift indicators in real time. You'll have visibility into error rates, response latency, and model confidence distributions. We also configure evaluation pipelines that run on a schedule and alert your team if performance degrades below agreed thresholds.

What's the difference between this and us just using ChatGPT ourselves?



ChatGPT is a general-purpose tool with no access to your internal data, no integration with your systems, no audit trail, and no reliability guarantees. Our deployment is purpose-built: it knows your company's specific data, integrates with your CRM, ERP, and documents, enforces your access controls, logs every interaction for compliance, and includes monitoring so performance issues are caught before they impact operations. The gap is between a consumer tool and an enterprise system.

Does your AI learn from our data and share it with other clients?



No. Each client deployment is fully isolated. Your data is never used to train shared models, and outputs from your deployment never influence another client's system. This is an architectural guarantee — there is no shared model layer across client deployments. The isolation is enforced by construction, not just by policy.

Do you have SOC 2 or ISO 27001 certification?



We do not currently hold these certifications, though our architecture is built with their requirements in mind. For enterprise procurement, we can provide a detailed security architecture document, data flow diagrams, and completed security questionnaire responses for your security team. Formal certification is on our roadmap — if this is a hard procurement requirement, discuss timing with your account contact so we can explore whether an architecture review satisfies your team in the interim.

What happens to our AI system if we stop working with us?



Because deployments live in your own cloud account or on-premises infrastructure, you retain full ownership of the system. The code, models, configuration, and data are all yours. We provide complete handover documentation and a transition process. You are not locked into our platform — we build systems you own, not subscriptions you rent.

Can the AI handle our industry-specific terminology?



Yes. During implementation we configure prompts and retrieval systems using the client's own internal documents, policy glossaries, and operational procedures. For lending clients, for example, the assistant understands LTV, DSR, HIBOR, and facility terms because it's trained on your materials — not generic internet text. Domain specificity is a core part of every implementation, not an add-on.

What level of support do we get after launch?



Post-launch support includes monitoring setup, a defined escalation path for performance issues, and scheduled review cycles to assess drift and retraining needs. Ongoing support scope depends on the engagement agreement — We offer both a managed service model (we monitor and maintain the system) and a handover model (client team owns operations with us on retainer for major model updates and retraining cycles).